

STA 314: Statistical Methods for Machine Learning I

Lecture 7 - Logistic regression and Gradient Descent

Xin Bing

Department of Statistical Sciences
University of Toronto

- In classification, $X \in \mathcal{X}$ and $Y \in C = \{0, 1, \dots, K - 1\}$.
- The Bayes rule

$$\arg \max_{k \in C} \mathbb{P} \{ Y = k \mid X = \mathbf{x} \}, \quad \forall \mathbf{x} \in \mathcal{X}$$

has the smallest expected error rate.

- For binary classification, our goal is to estimate

$$p(\mathbf{x}) = \mathbb{P} \{ Y = 1 \mid X = \mathbf{x} \}, \quad \forall \mathbf{x} \in \mathcal{X}.$$

Logistic Regression

Logistic Regression is a parametric approach that assumes parametric structure on

$$p(\mathbf{x}) = \mathbb{P}(Y = 1 \mid X = \mathbf{x}).$$

- It assumes

$$p(\mathbf{x}) = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p}}.$$

The function $f(t) = e^t / (1 + e^t)$ is called the logistic function.
 β_0, \dots, β_p are the parameters.

- It is easy to see that we always have $0 \leq p(\mathbf{x}) \leq 1$.
- Note that $p(\mathbf{x})$ is **NOT** a linear function either in \mathbf{x} or in $\beta = (\beta_0, \dots, \beta_p)$.

Logistic Regression

- A bit of rearrangement gives

$$\underbrace{\frac{p(\mathbf{x})}{1 - p(\mathbf{x})}}_{\text{odds}} = e^{\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p},$$
$$\underbrace{\log \left[\frac{p(\mathbf{x})}{1 - p(\mathbf{x})} \right]}_{\text{log-odds (a.k.a. logit)}} = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p.$$

odds $\in [0, \infty)$ and log-odds $\in (-\infty, \infty)$.

- Similar interpretation as linear models.
- How to estimate β ?

Maximum Likelihood Estimator (MLE)

Given $\mathcal{D}^{train} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$ with $y_i \in \{0, 1\}$, we estimate the parameters by **maximizing the likelihood** of \mathcal{D}^{train} .

The maximum likelihood principle

We seek the estimates of parameters such that the fitted probability are the closest to the individual's observed outcome.

Cont'd: MLE under logistic regression

Recipe of computing the MLE:

1. Write down the likelihood, as always!
2. Solve the optimization (maximization) problem.

The MLE has many nice properties!

- Asymp. consistent
- Asymp. normal
- Asymp. efficient

Inference under logistic regression

Let $\hat{\beta}$ be the MLE.

- Z-statistic is similar to t-statistic in regression, and is defined as

$$\frac{\hat{\beta}_j}{SE(\hat{\beta}_j)}, \quad \forall j \in \{0, 1, \dots, p\}.$$

- It produces p-value for testing the null hypothesis

$$H_0 : \beta_j = 0 \quad \text{v.s.} \quad H_1 : \beta_j \neq 0.$$

A large (absolute) value of the z-statistic or small p-value indicates evidence against H_0 .

Example: Default data

Consider the Default data using balance, income, and student status as predictors.

$$\log \left(\frac{p(X)}{1 - p(X)} \right) = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$$

$$p(X) = \frac{e^{\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p}}$$

	Coefficient	Std. Error	Z-statistic	P-value
Intercept	-10.8690	0.4923	-22.08	< 0.0001
balance	0.0057	0.0002	24.74	< 0.0001
income	0.0030	0.0082	0.37	0.7115
student [Yes]	-0.6468	0.2362	-2.74	0.0062

Prediction at **different levels** under logistic regression

Let $\hat{\boldsymbol{\beta}} = (\hat{\beta}_0, \dots, \hat{\beta}_p)$ be the MLE.

- Prediction of **the logit** at $\mathbf{x} \in \mathcal{X}$:

$$\text{logit}(\mathbf{x}) = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p.$$

- Prediction of **the conditional probability** $\mathbb{P}(Y = 1 \mid X = \mathbf{x})$:

$$\hat{\mathbb{P}}(Y = 1 \mid X = \mathbf{x}) = \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p}}$$

- Prediction of **the label** Y (i.e. *classification*) at $X = \mathbf{x}$:

$$\hat{y} = \begin{cases} 1, & \text{if } \hat{\mathbb{P}}(Y = 1 \mid X = \mathbf{x}) \geq 0.5; \\ 0, & \text{otherwise.} \end{cases}$$

Prediction of $\mathbb{P}(Y = 1 \mid X = \mathbf{x})$

Consider the Default data with student status as the only feature.

What is the probability of default for a student?

To fit the model, we encode student status as 1 for student and 0 otherwise.

	Coefficient	Std. Error	Z-statistic	P-value
Intercept	-3.5041	0.0707	-49.55	< 0.0001
student[Yes]	0.4049	0.1150	3.52	0.0004

$$\widehat{\Pr}(\text{default}=\text{Yes}|\text{student}=\text{Yes}) = \frac{e^{-3.5041+0.4049 \times 1}}{1 + e^{-3.5041+0.4049 \times 1}} = 0.0431,$$

$$\widehat{\Pr}(\text{default}=\text{Yes}|\text{student}=\text{No}) = \frac{e^{-3.5041+0.4049 \times 0}}{1 + e^{-3.5041+0.4049 \times 0}} = 0.0292.$$

Metrics used for evaluating classifiers

In classification, we have several metrics that can be used to evaluate a given classifier.

- The most commonly used metric is the overall classification accuracy.
- For binary classification, there are a few more out there.....

Logistic Regression on the Default Data

- Classify whether or not an individual will default on the basis of credit card balance and student status.
- The confusion matrix** of fitted logistic regression

		<i>True default status</i>		
		No	Yes	Total
<i>Predicted default status</i>	No	9,644	252	9,896
	Yes	23	81	104
	Total	9,667	333	10,000

Type of Errors for binary classification

		<i>True default status</i>		
		No	Yes	Total
<i>Predicted default status</i>	No	9,644	252	9,896
	Yes	23	81	104
	Total	9,667	333	10,000

1. The training error rate is $(23 + 252)/10000 = 2.75\%$.
2. **False positive rate (FPR)**: The fraction of negative examples that are classified as positive: $23/9667 = 0.2\%$ in default data.
3. **False negative rate (FNR)**: The fraction of positive examples that are classified as negative: $252/333 = 75.7\%$ in default data.¹

¹For a credit card company that is trying to identify high-risk individuals, the error rate 75.7% among individuals who default is unacceptable.

Types of Errors for binary classification

Q: How to modify the logistic classifier to lower the FNR?

- The current classifier is based on the rule

$$\hat{\mathbb{P}}(\text{default} = \text{yes} \mid X = \mathbf{x}) \geq 0.5.$$


- To lower FNR, we reduce the number of negative predictions.
Classify $X = \mathbf{x}$ to *yes* if

$$\hat{\mathbb{P}}(Y = \text{yes} \mid X = \mathbf{x}) \geq t.$$

for some $t < 0.5$.

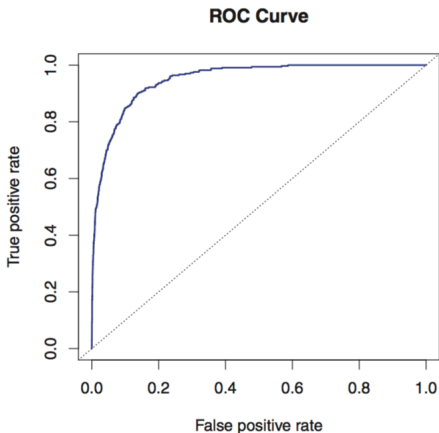
Trade-off between FPR and FNR

We can achieve better balance between FPR and FNR by varying the threshold:



ROC Curve

The **ROC curve** is a popular graphic for simultaneously displaying FPR and $\text{TPR} = 1 - \text{FNR}$ for all possible thresholds.



The overall performance of a classifier, summarized over all thresholds, is given by the area under the curve (**AUC**). High AUC is good.

More metrics in the binary classification

		<i>Predicted class</i>		
		– or Null	+ or Non-null	Total
<i>True class</i>	– or Null	True Neg. (TN)	False Pos. (FP)	N
	+ or Non-null	False Neg. (FN)	True Pos. (TP)	P
	Total	N^*	P^*	

Name	Definition	Synonyms
False Pos. rate	FP/N	Type I error, 1–Specificity
True Pos. rate	TP/P	1–Type II error, power, sensitivity, recall
Pos. Pred. value	TP/ P^*	Precision, 1–false discovery proportion
Neg. Pred. value	TN/ N^*	

The above also defines **sensitivity** and **specificity**.

Computation of the MLE under Logistic Regression

General steps of computing the MLE:

- Write down the likelihood, as always!
- Solve the optimization problem.

Likelihood under Logistic Regression

For simplicity, let us set $\beta_0 = 0$ such that

$$p(\mathbf{x}) = \frac{e^{\mathbf{x}^\top \boldsymbol{\beta}}}{1 + e^{\mathbf{x}^\top \boldsymbol{\beta}}}, \quad 1 - p(\mathbf{x}) = \frac{1}{1 + e^{\mathbf{x}^\top \boldsymbol{\beta}}}.$$

The data consists of $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ with

$$y_i \sim \text{Bernoulli}(p(\mathbf{x}_i)), \quad p(\mathbf{x}_i) = \frac{e^{\mathbf{x}_i^\top \boldsymbol{\beta}}}{1 + e^{\mathbf{x}_i^\top \boldsymbol{\beta}}}, \quad 1 \leq i \leq n.$$

- What is the likelihood of y_i ?

Likelihood under Logistic Regression

The likelihood of each data point (\mathbf{x}_i, y_i) at any β is

$$L(\beta; \mathbf{x}_i, y_i) = [p(\mathbf{x}_i)]^{y_i} [1 - p(\mathbf{x}_i)]^{1-y_i}$$

with

$$p(\mathbf{x}_i) = \frac{e^{\mathbf{x}_i^\top \beta}}{1 + e^{\mathbf{x}_i^\top \beta}}.$$

The joint likelihood of all data points is

$$L(\beta) = \prod_{i=1}^n [p(\mathbf{x}_i)]^{y_i} [1 - p(\mathbf{x}_i)]^{1-y_i}.$$

Log-likelihood under Logistic Regression

The log-likelihood at any β is

$$\begin{aligned}\ell(\beta) &= \log \left\{ \prod_{i=1}^n [p(\mathbf{x}_i)]^{y_i} [1 - p(\mathbf{x}_i)]^{1-y_i} \right\} \\&= \sum_{i=1}^n [y_i \log(p(\mathbf{x}_i)) + (1 - y_i) \log(1 - p(\mathbf{x}_i))] \\&= \sum_{i=1}^n \left[y_i \log \left(\frac{p(\mathbf{x}_i)}{1 - p(\mathbf{x}_i)} \right) + \log(1 - p(\mathbf{x}_i)) \right] \\&= \sum_{i=1}^n \left[y_i \mathbf{x}_i^\top \beta - \log \left(1 + e^{\mathbf{x}_i^\top \beta} \right) \right].\end{aligned}$$

How to compute the MLE?

How do we maximize the log-likelihood

$$\ell(\beta) = \sum_{i=1}^n \left[y_i \mathbf{x}_i^\top \beta - \log \left(1 + e^{\mathbf{x}_i^\top \beta} \right) \right]$$

for logistic regression?

- It is equivalent to minimize $-\ell(\beta)$ over β .
- No direct solution: taking derivatives of $\ell(\beta)$ w.r.t. β and setting them to 0 doesn't have an explicit solution.
- Need to use iterative procedure.

A general problem of solving a minimization problem

Suppose we want to solve the following problem

$$\hat{\mathbf{w}} = \operatorname{argmin}_{\mathbf{w} \in \Theta} \mathcal{J}(\mathbf{w}; \mathcal{D}^{train}) := \operatorname{argmin}_{\mathbf{w} \in \Theta} \mathcal{J}(\mathbf{w})$$

where $\mathcal{J}(\mathbf{w}; \mathcal{D}^{train})$ is a differentiable function in \mathbf{w} , and depends on \mathcal{D}^{train} as well, and Θ is a subspace of \mathbb{R}^p .

- The optimal solution (if exists) must be a **critical point**, i.e. point to which the derivative is zero (partial derivatives to zero for multi-dimensional parameter).

Finding the optimal solution requires to solve the equations

- **Partial derivatives:** derivatives of a multivariate function with respect to one of its arguments.

$$\frac{\partial}{\partial x_1} f(x_1, x_2) = \lim_{h \rightarrow 0} \frac{f(x_1 + h, x_2) - f(x_1, x_2)}{h}$$

- The minimum must occur at a point where the partial derivatives are zero.

$$\begin{bmatrix} \frac{\partial g}{\partial w_1} \\ \vdots \\ \frac{\partial g}{\partial w_p} \end{bmatrix} = 0$$

- This turns out to give a system of linear equations, which we can solve analytically in some scenarios.
- We may also use optimization techniques that iteratively get us closer to the solution.

Direct solution

- OLS:

$$\hat{\mathbf{w}} = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \mathcal{J}(\mathbf{w}) = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2.$$

The partial derivatives w.r.t. \mathbf{w} are

$$\frac{\partial g}{\partial \mathbf{w}} = -2\mathbf{X}^\top (\mathbf{y} - \mathbf{X}\mathbf{w}).$$

(If not familiar with multi-dimensional derivatives, calculate $\frac{\partial g}{\partial w_j}$ and stack them together).

Setting the above equal to zero results

$$\mathbf{X}^\top \mathbf{X} \hat{\mathbf{w}} = \mathbf{X}^\top \mathbf{y}, \quad \Rightarrow \quad \hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}.$$

- Ridge:

$$\hat{\mathbf{w}}_{\lambda}^R = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \mathcal{J}(\mathbf{w}) = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_2^2.$$

The partial derivatives w.r.t. \mathbf{w} are

$$\frac{\partial g}{\partial \mathbf{w}} = -2\mathbf{X}^{\top}(\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda\mathbf{w}.$$

Setting the above equal to zero results

$$(\mathbf{X}^{\top}\mathbf{X} + \lambda\mathbf{I}_p)\hat{\mathbf{w}}_{\lambda}^R = \mathbf{X}^{\top}\mathbf{y}, \quad \Rightarrow \quad \hat{\mathbf{w}}_{\lambda}^R = (\mathbf{X}^{\top}\mathbf{X} + \lambda\mathbf{I}_p)^{-1}\mathbf{X}^{\top}\mathbf{y}.$$

Gradient Descent

- Now let's see a second way to solve

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{J}(\mathbf{w})$$

which is more broadly applicable: **gradient descent**.



- Many times, we do not have a direct solution to

$$\frac{\partial \mathcal{J}}{\partial \mathbf{w}} = 0.$$

- Gradient descent is an **iterative algorithm**, which means we apply an update repeatedly until some criterion is met.

Gradient Descent

We **initialize** \mathbf{w} to something reasonable (e.g. all zeros) and repeatedly adjust them in the **direction of steepest descent**.



What is the direction of the steepest descent of $\mathcal{J}(\mathbf{w})$ at \mathbf{w} ?

Gradient Descent

- By definition, the direction of the greatest increase in $\mathcal{J}(\mathbf{w})$ at \mathbf{w} is its gradient $\partial\mathcal{J}/\partial\mathbf{w}$. So, we should update \mathbf{w} in the **opposite** direction of the gradient descent.
- The following update always decreases the cost function for small enough α (unless $\partial\mathcal{J}/\partial w_j = 0$): at the $(k + 1)$ th iteration,

$$w_j^{(k+1)} \leftarrow w_j^{(k)} - \alpha \cdot \frac{\partial\mathcal{J}}{\partial w_j} \Big|_{\mathbf{w}=\mathbf{w}^{(k)}}$$

- $\alpha > 0$ is a **learning rate** (or step size). The larger it is, the faster $\mathbf{w}^{(k+1)}$ changes relative to $\mathbf{w}^{(k)}$
 - ▶ We'll see later how to tune the learning rate, but values are typically small, e.g. 0.01 or 0.0001.

Example

$$\hat{\mathbf{w}} = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \mathcal{J}(\mathbf{w}), \quad \mathcal{J}(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2.$$

Update rule in vector form at the $k + 1$ th iteration:

$$\begin{aligned} \mathbf{w}^{(k+1)} &\leftarrow \mathbf{w}^{(k)} - \alpha \left. \frac{\partial \mathcal{J}}{\partial \mathbf{w}} \right|_{\mathbf{w}=\mathbf{w}^{(k)}} \\ &= \mathbf{w}^{(k)} + 2\alpha \mathbf{X}^\top (\mathbf{y} - \mathbf{X}\mathbf{w}^{(k)}). \end{aligned}$$

Initialization: $\mathbf{w}^{(0)} = \mathbf{0}$.

Stopping criteria

When do we stop?

- The objective value stops changing:

$$|\mathcal{J}(\mathbf{w}^{(k+1)}) - \mathcal{J}(\mathbf{w}^{(k)})| \text{ is small, i.e. } \leq 10^{-6}.$$

- The parameter stops changing: $\|\mathbf{w}^{(k+1)} - \mathbf{w}^{(k)}\|_2$ is small or $\|\mathbf{w}^{(k+1)} - \mathbf{w}^{(k)}\|_2 / \|\mathbf{w}^{(k)}\|_2$ is small.
- When we reach the maximum number (M) of iterations, e.g. $M = 1000$.

Gradient descent for solving the MLE under logistic regression

Recall we would like to solve

$$\min_{\mathbf{w} \in \mathbb{R}^p} \mathcal{J}(\mathbf{w})$$

where

$$\mathcal{J}(\mathbf{w}) = -\ell(\mathbf{w}) = \sum_{i=1}^n \left[-y_i \mathbf{x}_i^\top \mathbf{w} + \log \left(1 + e^{\mathbf{x}_i^\top \mathbf{w}} \right) \right].$$

The gradient at any \mathbf{w} is that, for any $j \in \{1, \dots, p\}$,

$$-\frac{\partial \ell(\mathbf{w})}{\partial w_j} = \sum_{i=1}^n \left[-y_i + \frac{e^{\mathbf{x}_i^\top \mathbf{w}}}{1 + e^{\mathbf{x}_i^\top \mathbf{w}}} \right] x_{ij} \quad (\text{verify this!})$$

Updates and stopping criteria

Therefore, at the $(k + 1)$ th iteration, with the learning rate α ,

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} - \alpha \sum_{i=1}^n \left[-y_i + \frac{e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}} \right] \mathbf{x}_i.$$

Initialization $\mathbf{w}^{(0)} = \mathbf{0}$.

- The objective value stops changing: $|\ell(\hat{\mathbf{w}}^{(k+1)}) - \ell(\hat{\mathbf{w}}^{(k)})|$ is small, say, $\leq 10^{-6}$.
- The parameter stops changing: $\|\hat{\mathbf{w}}^{(k+1)} - \hat{\mathbf{w}}^{(k)}\|_2$ is small or $\|\hat{\mathbf{w}}^{(k+1)} - \hat{\mathbf{w}}^{(k)}\|_2 / \|\hat{\mathbf{w}}^{(k)}\|_2$ is small.
- Stop after M iterations for some specified M , e.g. $M = 1000$.

When should we expect Gradient Descent to work?

Recall we try to solve

$$\hat{\mathbf{w}} = \underset{\mathbf{w} \in \Theta}{\operatorname{argmin}} \mathcal{J}(\mathbf{w}).$$

- Obviously, \mathcal{J} needs to be differentiable.
- If \mathcal{J} is also a convex function and Θ is a convex set, then Gradient Descent finds the optimal solution.
- In many cases, $\Theta = \mathbb{R}^p$ which is convex.

A set \mathcal{S} is convex if for any $\mathbf{x}_0, \mathbf{x}_1 \in \mathcal{S}$,

$$(1 - \lambda)\mathbf{x}_0 + \lambda\mathbf{x}_1 \in \mathcal{S} \quad \text{for all } 0 \leq \lambda \leq 1.$$


The Euclidean space \mathbb{R}^p is a convex set.

Convex Sets and Functions

- A function f is **convex** if for any $\mathbf{x}_0, \mathbf{x}_1$ in the domain of f ,

$$f((1 - \lambda)\mathbf{x}_0 + \lambda\mathbf{x}_1) \leq (1 - \lambda)f(\mathbf{x}_0) + \lambda f(\mathbf{x}_1), \quad \forall \lambda \in [0, 1].$$

- Equivalently, the set of points lying above the graph of f is convex.
- Intuitively: the function is bowl-shaped.



How to tell a loss is convex?

1. Verify the definition.
2. If f is twice differentiable and $f''(x) \geq 0$ for all x , then f is convex.
 - ▶ the least-squares loss function $(y - t)^2$ is convex as a function of t

- ▶ the function

$$-yt + \log(1 + e^t)$$

is convex in t .

3. There are other sufficient conditions for convex, but non-differentiable, functions!

4 A composition rule: linear functions preserve convexity.

- ▶ If f is a convex function and g is a linear function, then both $f \circ g$ and $g \circ f$ are convex.
 - ▶ the least-square loss $(y - \mathbf{x}^\top \mathbf{w})^2$ is convex in \mathbf{w}
 - ▶ the negative log-likelihood under logistic regression

$$-y\mathbf{x}^\top \mathbf{w} + \log\left(1 + e^{\mathbf{x}^\top \mathbf{w}}\right)$$

is convex in \mathbf{w} .

- ▶ Both $\sum_i (y_i - \mathbf{x}_i^\top \mathbf{w})^2$ and $\sum_i \left[-y_i \mathbf{x}_i^\top \mathbf{w} + \log\left(1 + e^{\mathbf{x}_i^\top \mathbf{w}}\right)\right]$ are convex in \mathbf{w} .

There are more composition rules!

A great book:

Convex Optimization, Stephen Boyd and Lieven Vandenberghe.

Gradient Descent for Linear Regression

- The squared error loss

$$\sum_{i=1} (y_i - \mathbf{x}_i^\top \mathbf{w})^2$$

of linear regression is a convex function. So there is a unique solution.

- Even in this case, we sometimes need to use GD.
- Why gradient descent, if we can find the optimum directly?
 - ▶ When p is large, GD is more efficient than direct solution
 - ▶ Linear regression solution: $(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$
 - ▶ Matrix inversion is an $\mathcal{O}(p^3)$ algorithm
 - ▶ Each GD update costs $\mathcal{O}(np)$
 - ▶ Or less with stochastic GD (Stochastic GD, later)
 - ▶ Huge difference if $p \gg \sqrt{n}$

Gradient descent for solving the MLE under logistic regression

- The negative log-likelihood


$$-\ell(\mathbf{w}) = \sum_{i=1}^n \left[-y_i \mathbf{x}_i^\top \mathbf{w} + \log \left(1 + e^{\mathbf{x}_i^\top \mathbf{w}} \right) \right]$$

is convex in \mathbf{w} .


- So we can use gradient descent to find the minima of the logistic loss!
- GD can be applied to more general settings!

Effect of the Learning Rate (Step Size)


- In gradient descent, the learning rate α is a hyperparameter we need to tune. Here are some things that can go wrong:



α too small:
slow progress



α too large:
oscillations




α much too large:
instability

- Good values are typically between 0.001 and 0.1. You should do a grid search if you want good performance (i.e. try 0.1, 0.03, 0.01, ...).

Training Curves

- To diagnose optimization problems, it's useful to look at the **training cost**: plot the training cost as a function of iteration.



- **Warning:** the training cost could be used to check whether the optimization problem reaches certain convergence. But
 - ▶ It does not tell whether we reach the global minimum or not
 - ▶ It does not tell anything on the performance of the fitted model

Gradient descent

Visualization:

http://www.cs.toronto.edu/~guerzhoy/321/lec/W01/linear_regression.pdf#page=21

Batch Gradient Descent

- Recall that

- OLS:

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \sum_{i=1}^n \left[y_i - \mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)} \right] \mathbf{x}_i.$$

- Logistic regression:

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \sum_{i=1}^n \left[y_i - \frac{e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}} \right] \mathbf{x}_i.$$

- Computing the gradient requires summing over **all** of the training examples, which can be done via matrix / vector operations. The fact that it uses all training samples is known as **batch training**.

Stochastic Gradient Descent

- Batch training is impractical if you have a large dataset (e.g. millions of training examples, $n \approx 10$ millions)!
- **Stochastic gradient descent (SGD)**: update the parameters based on the gradient for a single training example.

For each iteration $k \in \{1, 2, \dots\}$,

1. Choose $i \in \{1, \dots, n\}$ uniformly at random
2. Update the parameters by ONLY using this i th sample,

$$\begin{aligned}\hat{\mathbf{w}}^{(k+1)} &= \hat{\mathbf{w}}^{(k)} + \alpha \left[y_i - \mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)} \right] \mathbf{x}_i \\ \hat{\mathbf{w}}^{(k+1)} &= \hat{\mathbf{w}}^{(k)} + \alpha \left[y_i - \frac{e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}} \right] \mathbf{x}_i.\end{aligned}$$

Stochastic Gradient Descent

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \left[y_i - \mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)} \right] \mathbf{x}_i$$
$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \left[y_i - \frac{e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}} \right] \mathbf{x}_i.$$

Pros:

- Computational cost of each SGD update is independent of n !
- SGD can make significant progress before even seeing all the data!
- Mathematical justification: the gradients between SGD and GD have the same expectation for i.i.d. data.

Stochastic Gradient Descent

Cons: using single training example to estimate gradient:


- Variance in the estimate may be high

Compromise approach:


- compute the gradients on a randomly chosen medium-sized set of training examples $\mathcal{M} \subset \{1, \dots, n\}$, called a **mini-batch**.
- Stochastic gradients computed on larger mini-batches have smaller variance.
- The mini-batch size $|\mathcal{M}|$ is a hyperparameter that needs to be set.

Stochastic Gradient Descent

- Batch gradient descent moves directly downhill. SGD takes steps in a noisy direction, but moves downhill on average.




batch gradient descent




stochastic gradient descent

Learning Rate


- In gradient descent, the learning rate α is a hyperparameter we need to tune. Here are some things that can go wrong:



α too small:
slow progress



α too large:
oscillations




α much too large:
instability

- Good values are typically small. You should do a grid search if you want good performance (i.e. try 0.1, 0.03, 0.01, ...).


SGD Learning Rate

- In stochastic training, the learning rate also influences the **fluctuations** due to the stochasticity of the gradients.

small learning rate



large learning rate



- Typical strategy:
 - ▶ Use a large learning rate early in training so you can get close to the optimum
 - ▶ Gradually decay the learning rate to reduce the fluctuations