

# STA 314: Statistical Methods for Machine Learning I

## Lecture 8 - Logistic regression, gradient descent

Xin Bing

Department of Statistical Sciences  
University of Toronto

- In classification,  $X \in \mathbb{R}^p$  and  $Y \in C = \{0, 1, \dots, K - 1\}$ .
- The Bayes rule

$$\arg \max_{k \in C} \mathbb{P} \{ Y = k \mid X = x \}, \quad \forall x \in \mathbb{R}^p$$

has the smallest expected error rate.

- For binary classification, our goal is to estimate

$$p(x) = \mathbb{P} \{ Y = 1 \mid X = x \}, \quad \forall x \in \mathbb{R}^p.$$

# Logistic Regression

Logistic Regression is a parametric approach that assumes parametric structure on

$$p(X) = \mathbb{P}(Y = 1 \mid X).$$

- It assumes

$$p(X) = \frac{e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}}.$$

The function  $f(t) = e^t / (1 + e^t)$  is called the logistic function.  
 $\beta_0, \dots, \beta_p$  are the parameters.

- It is easy to see that we always have  $0 \leq p(X) \leq 1$ .
- Note that  $p(X)$  is **NOT** a linear function either in  $X$  or in  $\beta$ .

# Logistic Regression

- A bit of rearrangement gives

$$\underbrace{\frac{p(X)}{1 - p(X)}}_{\text{odds}} = e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p},$$
$$\underbrace{\log \left[ \frac{p(X)}{1 - p(X)} \right]}_{\text{log-odds (a.k.a. logit)}} = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p.$$

odds  $\in [0, \infty)$  and log-odds  $\in (-\infty, \infty)$ .

- Similar interpretation as linear models.
- How to estimate  $\beta_0, \dots, \beta_p$ ?

# Maximum Likelihood Estimator (MLE)

Given  $\mathcal{D}^{train} = \{(x_1, y_1), \dots, (x_n, y_n)\}$  with  $y_i \in \{0, 1\}$ , we estimate the parameters by **maximizing the likelihood** of  $\mathcal{D}^{train}$ .

## The maximum likelihood principle

The maximum likelihood principle is that we seek the estimates of parameters such that the fitted probability are the closest to the individual's observed outcome.

## Cont'd: MLE under logistic regression

General steps of computing the MLE:

- Write down the likelihood, as always!
- Solve the optimization (maximization) problem.

The MLE has many nice properties!

- Asymp consistent.
- Asymp normal.
- And more.....

# Inference under logistic regression

Let  $\hat{\beta}$  be the MLE.

- Z-statistic is similar to t-statistic in regression, and is defined as

$$\frac{\hat{\beta}_j}{SE(\hat{\beta}_j)}, \quad \forall j \in \{0, 1, \dots, p\}.$$

- It produces p-value for testing the null hypothesis

$$H_0 : \beta_j = 0 \quad \text{v.s.} \quad H_1 : \beta_j \neq 0.$$

A large (absolute) value of the z-statistic or small p-value indicates evidence against  $H_0$ .

## Example: Default data

Consider the Default data using balance, income, and student status as predictors.

$$\log \left( \frac{p(X)}{1 - p(X)} \right) = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$$

$$p(X) = \frac{e^{\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p}}$$

	Coefficient	Std. Error	Z-statistic	P-value
Intercept	-10.8690	0.4923	-22.08	< 0.0001
balance	0.0057	0.0002	24.74	< 0.0001
income	0.0030	0.0082	0.37	0.7115
student [Yes]	-0.6468	0.2362	-2.74	0.0062



# Prediction at **different levels** under logistic regression

Let  $\hat{\beta}_0, \dots, \hat{\beta}_p$  be the MLE.

- Prediction of the logit at  $x \in \mathbb{R}^p$ :

$$\text{logit}(x) = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p.$$

- Prediction of  $\mathbb{P}(Y = 1 \mid X = x)$ :

$$\hat{\mathbb{P}}(Y = 1 \mid X = x) = \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p}}$$

- Prediction of  $Y$  (i.e. *classification*) at  $X = x$ :

$$\hat{y} = \begin{cases} 1, & \text{if } \hat{\mathbb{P}}(Y = 1 \mid X = x) \geq 0.5; \\ 0, & \text{otherwise.} \end{cases}$$

# Prediction of $\mathbb{P}(Y = 1 \mid X)$

Consider the Default data with student status as the only feature. What is our estimated probability of default for a student?

To fit the model, we encode student status as 1 for student and 0 otherwise.

	Coefficient	Std. Error	Z-statistic	P-value
<b>Intercept</b>	-3.5041	0.0707	-49.55	< 0.0001
<b>student[Yes]</b>	0.4049	0.1150	3.52	0.0004

$$\widehat{\Pr}(\text{default}=\text{Yes}|\text{student}=\text{Yes}) = \frac{e^{-3.5041+0.4049 \times 1}}{1 + e^{-3.5041+0.4049 \times 1}} = 0.0431,$$

$$\widehat{\Pr}(\text{default}=\text{Yes}|\text{student}=\text{No}) = \frac{e^{-3.5041+0.4049 \times 0}}{1 + e^{-3.5041+0.4049 \times 0}} = 0.0292.$$

# Metrics used for evaluating classifiers

In classification, we have several metrics that can be used to evaluate a given classifier.

- The most commonly used metric is the overall classification accuracy.
- For binary classification, there are a few more out there.....

# Logistic Regression on the Default Data

Classify whether or not an individual will default on the basis of credit card balance and student status. The confusion matrix on default data.

		<i>True default status</i>		
		No	Yes	Total
<i>Predicted default status</i>	No	9,644	252	9,896
	Yes	23	81	104
Total		9,667	333	10,000

- The training error rate is  $(23 + 252)/10000 = 2.75\%$ .
- **False positive rate (FPR)**: The fraction of negative examples that are classified as positive:  $23/9667 = 0.2\%$  in default data.
- **False negative rate (FNR)**: The fraction of positive examples that are classified as negative:  $252/333 = 75.7\%$  in default data.<sup>1</sup>

<sup>1</sup>For a credit card company that is trying to identify high-risk individuals, the error rate 75.7% among individuals who default is unacceptable.

# Types of Errors for binary classification

- The FNR is too high. How to modify the logistic classifier to lower the FNR?
- The current classifier is based on the rule

$$\hat{\mathbb{P}}(\text{default} = \text{yes} \mid X = x) \geq 0.5.$$


- To lower FNR, we reduce the number of negative predictions. Classify  $X = x$  to yes if

$$\hat{\mathbb{P}}(Y = \text{yes} \mid X = x) \geq t.$$

for some  $t < 0.5$ .


# Trade-off between FPR and FNR

We can achieve better balance between FPR and FNR by varying the threshold:



# ROC Curve

The **ROC curve** is a popular graphic for simultaneously displaying FPR and TPR for all possible thresholds.



The overall performance of a classifier, summarized over all thresholds, is given by the area under the curve (**AUC**). High AUC is good.

# More metrics in the binary classification

		<i>Predicted class</i>		
		– or Null	+ or Non-null	Total
<i>True class</i>	– or Null	True Neg. (TN)	False Pos. (FP)	N
	+ or Non-null	False Neg. (FN)	True Pos. (TP)	P
	Total	N*	P*	

Name	Definition	Synonyms
False Pos. rate	FP/N	Type I error, 1–Specificity
True Pos. rate	TP/P	1–Type II error, power, sensitivity, recall
Pos. Pred. value	TP/P*	Precision, 1–false discovery proportion
Neg. Pred. value	TN/N*	

The above also defines **sensitivity** and **specificity**.



# Computation of the MLE under Logistic Regression

General steps of computing the MLE:

- Write down the likelihood, as always!
- Solve the optimization problem.

# Likelihood under Logistic Regression

For simplicity, let us set  $\beta_0 = 0$  such that

$$p(x) = \frac{e^{x^\top \beta}}{1 + e^{x^\top \beta}}, \quad 1 - p(x) = \frac{1}{1 + e^{x^\top \beta}}.$$

The data consists of  $(x_1, y_1), \dots, (x_n, y_n)$  with

$$y_i \sim \text{Bernoulli}(p(x_i)), \quad p(x_i) = \frac{e^{x_i^\top \beta}}{1 + e^{x_i^\top \beta}}, \quad 1 \leq i \leq n.$$

- What is the likelihood of  $y_i$ ?

# Likelihood under Logistic Regression

The likelihood of each data point  $(x_i, y_i)$  at any  $\beta$  is

$$L_i(\beta) = [p(x_i)]^{y_i} [1 - p(x_i)]^{1-y_i}$$

with

$$p(x_i) = \frac{e^{x_i^\top \beta}}{1 + e^{x_i^\top \beta}}.$$

The joint likelihood of all data points is

$$L(\beta) = \prod_{i=1}^n [p(x_i)]^{y_i} [1 - p(x_i)]^{1-y_i}.$$

# Log-likelihood under Logistic Regression

The log-likelihood at any  $\beta$  is

$$\begin{aligned}\ell(\beta) &= \log \left\{ \prod_{i=1}^n [p(x_i)]^{y_i} [1 - p(x_i)]^{1-y_i} \right\} \\&= \sum_{i=1}^n [y_i \log(p(x_i)) + (1 - y_i) \log(1 - p(x_i))] \\&= \sum_{i=1}^n \left[ y_i \log \left( \frac{p(x_i)}{1 - p(x_i)} \right) + \log(1 - p(x_i)) \right] \\&= \sum_{i=1}^n \left[ y_i x_i^\top \beta - \log \left( 1 + e^{x_i^\top \beta} \right) \right].\end{aligned}$$

# How to compute the MLE?

How do we maximize the log-likelihood

$$\ell(\beta) = \sum_{i=1}^n \left[ y_i x_i^\top \beta - \log \left( 1 + e^{x_i^\top \beta} \right) \right]$$

for logistic regression?

- It is equivalent to minimize  $-\ell(\beta)$  over  $\beta$ .
- No direct solution: taking derivatives of  $\ell(\beta)$  w.r.t.  $\beta$  and setting them to 0 doesn't have an explicit solution.
- Need to use iterative procedure.

# A general problem of solving a minimization problem

Suppose we want to solve the following problem

$$\hat{\mathbf{w}} = \operatorname{argmin}_{\mathbf{w} \in \Theta} \mathcal{J}(\mathbf{w}; \mathcal{D}^{train}) := \operatorname{argmin}_{\mathbf{w} \in \Theta} \mathcal{J}(\mathbf{w})$$

where  $\mathcal{J}(\mathbf{w}; \mathcal{D}^{train})$  is a differentiable function in  $\mathbf{w}$ , and depends on  $\mathcal{D}^{train}$  as well, and  $\Theta$  is a subspace of  $\mathbb{R}^p$ .

- The optimal solution (if exists) must be a **critical point**, i.e. point to which the derivative is zero (partial derivatives to zero for multi-dimensional parameter).

# Finding the optimal solution requires to solve the equations

- **Partial derivatives:** derivatives of a multivariate function with respect to one of its arguments.

$$\frac{\partial}{\partial x_1} f(x_1, x_2) = \lim_{h \rightarrow 0} \frac{f(x_1 + h, x_2) - f(x_1, x_2)}{h}$$

- The minimum must occur at a point where the partial derivatives are zero.

$$\begin{bmatrix} \frac{\partial g}{\partial w_1} \\ \vdots \\ \frac{\partial g}{\partial w_p} \end{bmatrix} = 0$$

- This turns out to give a system of linear equations, which we can solve analytically in some scenarios.
- We may also use optimization techniques that iteratively get us closer to the solution.

# Direct solution

- OLS:

$$\hat{\mathbf{w}} = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \mathcal{J}(\mathbf{w}) = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2.$$

The partial derivatives w.r.t.  $\mathbf{w}$  are

$$\frac{\partial g}{\partial \mathbf{w}} = -2\mathbf{X}^\top (\mathbf{y} - \mathbf{X}\mathbf{w}).$$

(If not familiar with multi-dimensional derivatives, calculate  $\frac{\partial g}{\partial w_j}$  and stack them together).

Setting the above equal to zero results

$$\mathbf{X}^\top \mathbf{X} \hat{\mathbf{w}} = \mathbf{X}^\top \mathbf{y}, \quad \Rightarrow \quad \hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}.$$



- Ridge:

$$\hat{\mathbf{w}}_{\lambda}^R = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \mathcal{J}(\mathbf{w}) = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_2^2.$$

The partial derivatives w.r.t.  $\mathbf{w}$  are

$$\frac{\partial g}{\partial \mathbf{w}} = -2\mathbf{X}^{\top}(\mathbf{y} - \mathbf{X}\mathbf{w}) + 2\lambda\mathbf{w}.$$

Setting the above equal to zero results

$$(\mathbf{X}^{\top}\mathbf{X} + \lambda\mathbf{I}_p)\hat{\mathbf{w}}_{\lambda}^R = \mathbf{X}^{\top}\mathbf{y}, \quad \Rightarrow \quad \hat{\mathbf{w}}_{\lambda}^R = (\mathbf{X}^{\top}\mathbf{X} + \lambda\mathbf{I}_p)^{-1}\mathbf{X}^{\top}\mathbf{y}.$$

- Now let's see a second way to solve

$$\hat{\mathbf{w}} = \min_{\mathbf{w}} \mathcal{J}(\mathbf{w})$$

which is more broadly applicable: [gradient descent](#).



- Many times, we do not have a direct solution to

$$\frac{\partial \mathcal{J}}{\partial \mathbf{w}} = 0.$$

- Gradient descent is an [iterative algorithm](#), which means we apply an update repeatedly until some criterion is met.

# Gradient Descent

We **initialize**  $\mathbf{w}$  to something reasonable (e.g. all zeros) and repeatedly adjust them in the **direction of steepest descent**.



What is the direction of the steepest descent of  $\mathcal{J}(\mathbf{w})$  at  $\mathbf{w}$ ?

# Gradient Descent

- By definition, the direction of the greatest increase in  $\mathcal{J}(\mathbf{w})$  at  $\mathbf{w}$  is its gradient  $\partial\mathcal{J}/\partial\mathbf{w}$ . So, we should update  $\mathbf{w}$  in the **opposite** direction of the gradient descent.
- The following update always decreases the cost function for small enough  $\alpha$  (unless  $\partial\mathcal{J}/\partial w_j = 0$ ): at the  $(k + 1)$ th iteration,

$$w_j^{(k+1)} \leftarrow w_j^{(k)} - \alpha \frac{\partial \mathcal{J}}{\partial w_j} \Big|_{\mathbf{w}=\mathbf{w}^{(k)}}$$

- $\alpha > 0$  is a **learning rate** (or step size). The larger it is, the faster  $\mathbf{w}^{(k+1)}$  changes relative to  $\mathbf{w}^{(k)}$ 
  - ▶ We'll see later how to tune the learning rate, but values are typically small, e.g. 0.01 or 0.0001.

## Example

$$\hat{\mathbf{w}} = \underset{\mathbf{w} \in \mathbb{R}^p}{\operatorname{argmin}} \mathcal{J}(\mathbf{w}), \quad \mathcal{J}(\mathbf{w}) = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2.$$

Update rule in vector form at the  $k + 1$ th iteration:

$$\begin{aligned} \mathbf{w}^{(k+1)} &\leftarrow \mathbf{w}^{(k)} - \alpha \left. \frac{\partial \mathcal{J}}{\partial \mathbf{w}} \right|_{\mathbf{w}=\mathbf{w}^{(k)}} \\ &= \mathbf{w}^{(k)} + 2\alpha \mathbf{X}^\top (\mathbf{y} - \mathbf{X}\mathbf{w}^{(k)}). \end{aligned}$$

Initialization:  $\mathbf{w}^{(0)} = \mathbf{0}$ .

# Stopping criteria

When do we stop?

- The objective value stops changing:

$$|\mathcal{J}(\mathbf{w}^{(k+1)}) - \mathcal{J}(\mathbf{w}^{(k)})| \text{ is small, i.e. } \leq 10^{-6}.$$

- The parameter stops changing:  $\|\mathbf{w}^{(k+1)} - \mathbf{w}^{(k)}\|_2$  is small or  $\|\mathbf{w}^{(k+1)} - \mathbf{w}^{(k)}\|_2 / \|\mathbf{w}^{(k)}\|_2$  is small.
- When we reach the maximum number ( $M$ ) of iterations, e.g.  $M = 1000$ .

# Gradient descent for solving the MLE under logistic regression

Recall we would like to solve

$$\min_{\mathbf{w} \in \mathbb{R}^p} \mathcal{J}(\mathbf{w})$$

where

$$\mathcal{J}(\mathbf{w}) = -\ell(\mathbf{w}) = \sum_{i=1}^n \left[ -y_i x_i^\top \mathbf{w} + \log \left( 1 + e^{x_i^\top \mathbf{w}} \right) \right].$$

The gradient at any  $\mathbf{w}$  is that, for any  $j \in \{1, \dots, p\}$ ,

$$\frac{\partial [-\ell(\mathbf{w})]}{\partial \beta_j} = \sum_{i=1}^n \left[ -y_i + \frac{e^{x_i^\top \mathbf{w}}}{1 + e^{x_i^\top \mathbf{w}}} \right] x_{ij} \quad (\text{verify this!})$$

# Updates and stopping criteria

Therefore, at the  $(k + 1)$ th iteration, with the learning rate  $\alpha$ ,

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} - \alpha \sum_{i=1}^n \left[ -y_i + \frac{e^{x_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{x_i^\top \hat{\mathbf{w}}^{(k)}}} \right] x_i.$$

Initialization  $\mathbf{w}^{(0)} = \mathbf{0}$ .

- The objective value stops changing:  $|\ell(\hat{\mathbf{w}}^{(k+1)}) - \ell(\hat{\mathbf{w}}^{(k)})|$  is small, say,  $\leq 10^{-6}$ .
- The parameter stops changing:  $\|\hat{\mathbf{w}}^{(k+1)} - \hat{\mathbf{w}}^{(k)}\|_2$  is small or  $\|\hat{\mathbf{w}}^{(k+1)} - \hat{\mathbf{w}}^{(k)}\|_2 / \|\hat{\mathbf{w}}^{(k)}\|_2$  is small.
- Stop after  $M$  iterations for some specified  $M$ , e.g.  $M = 1000$ .



# When should we expect Gradient Descent to work?

Recall we try to solve

$$\hat{\mathbf{w}} = \min_{\mathbf{w} \in \Theta} \mathcal{J}(\mathbf{w}).$$

- Obviously,  $\mathcal{J}$  needs to be differentiable.
- If  $\mathcal{J}$  is also a convex function and  $\Theta$  is a convex set, then Gradient Descent finds the optimal solution.
- In many cases,  $\Theta = \mathbb{R}^p$  which is convex.

A set  $\mathcal{S}$  is convex if for any  $\mathbf{x}_0, \mathbf{x}_1 \in \mathcal{S}$ ,

$$(1 - \lambda)\mathbf{x}_0 + \lambda\mathbf{x}_1 \in \mathcal{S} \quad \text{for all } 0 \leq \lambda \leq 1.$$


The Euclidean space  $\mathbb{R}^p$  is a convex set.

# Convex Sets and Functions

- A function  $f$  is **convex** if for any  $\mathbf{x}_0, \mathbf{x}_1$  in the domain of  $f$ ,

$$f((1 - \lambda)\mathbf{x}_0 + \lambda\mathbf{x}_1) \leq (1 - \lambda)f(\mathbf{x}_0) + \lambda f(\mathbf{x}_1), \quad \forall \lambda \in [0, 1].$$

- Equivalently, the set of points lying above the graph of  $f$  is convex.
- Intuitively: the function is bowl-shaped.



# How to tell a loss is convex?

1. Verify the definition.
2. If  $f$  is twice differentiable and  $f''(x) \geq 0$  for all  $x$ , then  $f$  is convex.
  - ▶ the least-squares loss function  $(y - t)^2$  is convex as a function of  $t$
  - ▶ the function
$$-yt + \log(1 + e^t)$$
is convex in  $t$ .
3. There are other sufficient conditions for convex, but non-differentiable, functions!

#### 4 A composition rule: linear functions preserve convexity.

- ▶ If  $f$  is a convex function and  $g$  is a linear function, then both  $f \circ g$  and  $g \circ f$  are convex.
  - ▶ the least-square loss  $(y - x^\top \mathbf{w})^2$  is convex in  $\mathbf{w}$
  - ▶ the negative log-likelihood under logistic regression

$$-y x^\top \mathbf{w} + \log \left( 1 + e^{x^\top \mathbf{w}} \right)$$

is convex in  $\mathbf{w}$ .

- ▶ Both  $\sum_i (y_i - x_i^\top \mathbf{w})^2$  and  $\sum_i \left[ -y_i x_i^\top \mathbf{w} + \log \left( 1 + e^{x_i^\top \mathbf{w}} \right) \right]$  are convex in  $\mathbf{w}$ .

There are more composition rules!

A great book:

*Convex Optimization, Stephen Boyd and Lieven Vandenberghe.*

# Gradient Descent for Linear Regression

- The squared error loss

$$\sum_{i=1} (y_i - \mathbf{x}_i^T \mathbf{w})^2$$

of linear regression is a convex function. So there is a unique solution. Even in this case, we sometimes need to use GD.

- Why gradient descent, if we can find the optimum directly?
  - ▶ When  $p$  is large, GD is more efficient than direct solution
    - ▶ Linear regression solution:  $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
    - ▶ Matrix inversion is an  $\mathcal{O}(p^3)$  algorithm
    - ▶ Each GD update costs  $\mathcal{O}(np)$
    - ▶ Or less with stochastic GD (Stochastic GD, later)
    - ▶ Huge difference if  $p \gg \sqrt{n}$

# Gradient descent for solving the MLE under logistic regression

- The negative log-likelihood


$$-\ell(\mathbf{w}) = \sum_{i=1}^n \left[ -y_i x_i^\top \mathbf{w} + \log \left( 1 + e^{x_i^\top \mathbf{w}} \right) \right]$$

is convex in  $\mathbf{w}$ .


- So we can use gradient descent to find the minima of the logistic loss!
- GD can be applied to more general settings!

# Effect of the Learning Rate (Step Size)


- In gradient descent, the learning rate  $\alpha$  is a hyperparameter we need to tune. Here are some things that can go wrong:



$\alpha$  too small:  
slow progress



$\alpha$  too large:  
oscillations




$\alpha$  much too large:  
instability

- Good values are typically between 0.001 and 0.1. You should do a grid search if you want good performance (i.e. try 0.1, 0.03, 0.01, ...).



# Training Curves

- To diagnose optimization problems, it's useful to look at the **training cost**: plot the training cost as a function of iteration.



- **Warning:** the training cost could be used to check whether the optimization problem reaches certain convergence. But
  - ▶ It does not tell whether we reach the global minimum or not
  - ▶ It does not tell anything on the performance of the fitted model

# Gradient descent

Visualization:

[http://www.cs.toronto.edu/~guerzhoy/321/lec/W01/linear\\_regression.pdf#page=21](http://www.cs.toronto.edu/~guerzhoy/321/lec/W01/linear_regression.pdf#page=21)

# Batch Gradient Descent

- Recall that

- ▶ OLS:

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \sum_{i=1}^n \left[ y_i - x_i^\top \hat{\mathbf{w}}^{(k)} \right] x_i.$$

- ▶ Logistic regression:

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \sum_{i=1}^n \left[ y_i - \frac{e^{x_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{x_i^\top \hat{\mathbf{w}}^{(k)}}} \right] x_i.$$

- Computing the gradient requires summing over **all** of the training examples, which can be done via matrix / vector operations. The fact that it uses all training samples is known as **batch training**.

# Stochastic Gradient Descent

- Batch training is impractical if you have a large dataset (e.g. millions of training examples,  $n \approx 10$  millions)!
- **Stochastic gradient descent (SGD)**: update the parameters based on the gradient for a single training example.

For each iteration  $k \in \{1, 2, \dots\}$ ,

1. Choose  $i \in \{1, \dots, n\}$  uniformly at random
2. Update the parameters by **ONLY** using this  $i$ th sample,

$$\begin{aligned}\hat{\mathbf{w}}^{(k+1)} &= \hat{\mathbf{w}}^{(k)} + \alpha \left[ y_i - \mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)} \right] \mathbf{x}_i \\ \hat{\mathbf{w}}^{(k+1)} &= \hat{\mathbf{w}}^{(k)} + \alpha \left[ y_i - \frac{e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{\mathbf{x}_i^\top \hat{\mathbf{w}}^{(k)}}} \right] \mathbf{x}_i.\end{aligned}$$

# Stochastic Gradient Descent

$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \left[ y_i - x_i^\top \hat{\mathbf{w}}^{(k)} \right] x_i$$
$$\hat{\mathbf{w}}^{(k+1)} = \hat{\mathbf{w}}^{(k)} + \alpha \left[ y_i - \frac{e^{x_i^\top \hat{\mathbf{w}}^{(k)}}}{1 + e^{x_i^\top \hat{\mathbf{w}}^{(k)}}} \right] x_i.$$

## Pros:

- Computational cost of each SGD update is independent of  $n$ !
- SGD can make significant progress before even seeing all the data!
- Mathematical justification: the gradients between SGD and GD have the same expectation for i.i.d. data.

# Stochastic Gradient Descent

**Cons:** using single training example to estimate gradient:


- Variance in the estimate may be high

Compromise approach:


- compute the gradients on a randomly chosen medium-sized set of training examples  $\mathcal{M} \subset \{1, \dots, n\}$ , called a **mini-batch**.
- Stochastic gradients computed on larger mini-batches have smaller variance.
- The mini-batch size  $|\mathcal{M}|$  is a hyperparameter that needs to be set.

# Stochastic Gradient Descent

- Batch gradient descent moves directly downhill. SGD takes steps in a noisy direction, but moves downhill on average.




**batch gradient descent**




**stochastic gradient descent**

# Learning Rate


- In gradient descent, the learning rate  $\alpha$  is a hyperparameter we need to tune. Here are some things that can go wrong:



$\alpha$  too small:  
slow progress



$\alpha$  too large:  
oscillations



$\alpha$  much too large:  
instability


- Good values are typically small. You should do a grid search if you want good performance (i.e. try 0.1, 0.03, 0.01, ...).




# SGD Learning Rate

- In stochastic training, the learning rate also influences the **fluctuations** due to the stochasticity of the gradients.

small learning rate



large learning rate



- Typical strategy:
  - ▶ Use a large learning rate early in training so you can get close to the optimum
  - ▶ Gradually decay the learning rate to reduce the fluctuations