

# STA 314: Statistical Methods for Machine Learning I

## Lecture 9 - Moving beyond linearity, $k$ -Nearest Neighbor

Xin Bing

Department of Statistical Sciences  
University of Toronto

# Moving Beyond Linearity

The linearity assumption in the feature space is almost always an approximation, and sometimes a poor one.

We consider the following extensions to relax the linearity assumption.

- Univariate feature ( $p = 1$ ):
  - ▶ Polynomial regression
  - ▶ Step functions
  - ▶ Regression splines
- Multivariate feature ( $p > 1$ ):
  - ▶ Local regression
  - ▶ Generalized additive models

# Polynomial Regression

- The **polynomial regression**

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \cdots + \beta_d x_i^d + \epsilon_i,$$

where  $\epsilon_i$  is the error term and  $x_i \in \mathbb{R}$ .

- Can be fitted by the OLS approach.
- Coefficients themselves are not interpretable; we are more interested in the trend of the fitted function

$$\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0 + \hat{\beta}_2 x_0^2 + \cdots + \hat{\beta}_d x_0^d.$$

- The degree  $d$  in practice is typically no greater than 4, and can be chosen via cross-validation.

# Polynomial Regression

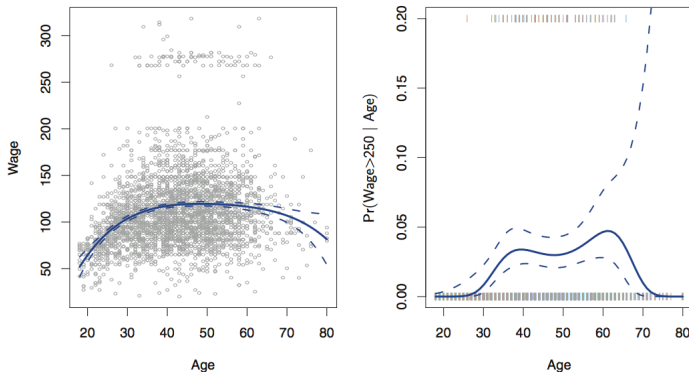
- The polynomial regression can be used for classification as well.
  - ▶ For instance, in the logistic regression,

$$\text{logit}(\mathbb{P}(Y_i = 1 \mid X_i = x_i)) = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \cdots + \beta_d x_i^d.$$

- ▶ Can be fit by maximizing the likelihood.
- However, polynomials have notorious tail behavior – very bad for extrapolation.

# The Wage Data

Degree-4 Polynomial



Left: The solid blue curve is a degree-4 polynomial of wage as a function of age, fit by the OLS. The dotted curves are estimated 95 % confidence intervals.

Right: We model the binary event  $1\{\text{wage} > 250\}$  using logistic regression, with a degree-4 polynomial.

# Step Functions

- The polynomial regression imposes a global structure on the non-linearity of  $X$ .
- The **step function** approach avoids such a global structure by breaking the range of  $X$  into bins.
- For pre-specified  $K$  cut points  $c_1, \dots, c_K$ , define

$$\begin{aligned}C_0(X) &= 1\{X < c_1\}, \\C_1(X) &= 1\{c_1 \leq X < c_2\}, \\&\vdots \\C_K(X) &= 1\{c_K \leq X\}.\end{aligned}$$

$C_0(X), \dots, C_K(X)$  are in fact  $(K + 1)$  dummy variables, and they sum up to 1.

# Step Functions

- Step function approach assumes


$$y_i = \beta_0 + \beta_1 C_1(x_i) + \beta_2 C_2(x_i) + \dots + \beta_K C_K(x_i) + \epsilon_i,$$

where  $\epsilon_i$  is the error term. (Note we don't need  $C_0(x_i)$  in the model when we also have the intercept term  $\beta_0$ .)

- Can be fitted by the OLS.
- $\beta_j$  represents the average change in the response  $Y$  for  $c_j \leq X < c_{j+1}$  relative to  $X < c_1$ .

# The Wage Data

## Piecewise Constant



Left: The solid blue curve is a step function of wage as a function of age, fit by least squares. The dotted curves indicate an estimated 95 % confidence interval.

Right: We model the binary event  $\text{wage} > 250$  using logistic regression, with the step function.



# Pros and Cons of Step Function

- The step function approach is widely used in biostatistics and epidemiology among other areas, because the model is easy to fit and the regression coefficient has a natural interpretation.
- However, piecewise-constant functions can miss the trend of the true relationship between  $Y$  and  $X$ . The choice of cut points can be difficult to specify.
- How about combining polynomial and step function?

# Piecewise Polynomials


- Instead of a single polynomial in  $X$  over its whole domain, we can use different polynomials in different regions:

$$y_i = \begin{cases} \beta_{01} + \beta_{11}x_i + \beta_{21}x_i^2 + \beta_{31}x_i^3 + \epsilon_i & \text{if } x_i < c; \\ \beta_{02} + \beta_{12}x_i + \beta_{22}x_i^2 + \beta_{32}x_i^3 + \epsilon_i & \text{if } x_i \geq c. \end{cases}$$


- The cut point  $c$  is called **knot**. Using more knots leads to a more flexible piecewise polynomial.
- In general, if we place  $K$  different knots throughout the range of  $X$ , then we will end up fitting  $(K + 1)$  different cubic polynomials.

# The Wage Data


**Piecewise Cubic**




**Continuous Piecewise Cubic**



**Cubic Spline**



**Linear Spline**



- Better to add constraints to polynomials at the knots for:
  - ▶ continuity: equal function values
  - ▶ smoothness: equal first and second order derivatives
  - ▶ higher order derivatives
- The constrained polynomials are called **splines**. A degree- $d$  spline contains piecewise degree- $d$  polynomials, with continuity in derivatives up to degree  $(d - 1)$  at each knot.
- How can we construct the degree- $d$  spline?

# Linear Splines

- A **linear spline** has piecewise linear functions continuous at each knot. That is, with knots at  $\xi_1 < \xi_2 < \dots < \xi_K$ ,

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 (x_i - \xi_1)_+ \dots + \beta_{K+1} (x_i - \xi_K)_+ + \epsilon_i,$$

where

$$(x_i - \xi_k)_+ = \begin{cases} x_i - \xi_k & \text{if } x_i > \xi_k \\ 0 & \text{otherwise} \end{cases}.$$

- A basis representation:

$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+1} b_{K+1}(x_i) + \epsilon_i,$$

where  $b_k$  are **basis functions**

$$b_1(x_i) = x_i, \quad b_{k+1}(x_i) = (x_i - \xi_k)_+, \quad k = 1, \dots, K,$$

- Interpretation of  $\beta_1$ : the averaged increase of  $Y$  associated with one unit of  $X$  for  $X < \xi_1$ .

- A **cubic spline** has piecewise cubic polynomials with continuous derivatives up to order 2 at each knot. That is, with  $K$  knots at  $\xi_1 < \xi_2 < \dots < \xi_K$ ,


$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+3} b_{K+3}(x_i) + \epsilon_i,$$

where  $b_k$  are basis functions

$$b_1(x_i) = x_i, \quad b_2(x_i) = x_i^2, \quad b_3(x_i) = x_i^3,$$


$$b_{k+3}(x_i) = (x_i - \xi_k)_+^3, \quad k = 1, \dots, K.$$

# Cubic Splines



# Natural Splines

A natural spline is a regression spline with additional boundary constraints: the function is required to be linear at the boundary.





- Choosing the number and locations of the knots
  - ▶ Typically, we place  $K$  knots at certain quantiles of the data or place on the range of  $X$  with equal space. Oftentimes, the placement of knots is not very crucial.
  - ▶ We use cross-validation to choose  $K$ .
- Polynomial regressions and step functions are special cases of splines.
- Another variant: smoothing spline (ISLR 7.5).

# Local Regression

**Local regression** predicts at a target point  $x_0$  using only the nearby training observations.

---

**Algorithm 7.1** *Local Regression At  $X = x_0$* 

---


1. Gather the fraction  $s = k/n$  of training points whose  $x_i$  are closest to  $x_0$ .
2. Assign a weight  $K_{i0} = K(x_i, x_0)$  to each point in this neighborhood, so that the point furthest from  $x_0$  has weight zero, and the closest has the highest weight. All but these  $k$  nearest neighbors get weight zero.
3. Fit a *weighted least squares regression* of the  $y_i$  on the  $x_i$  using the aforementioned weights, by finding  $\hat{\beta}_0$  and  $\hat{\beta}_1$  that minimize

$$\sum_{i=1}^n K_{i0} (y_i - \beta_0 - \beta_1 x_i)^2. \quad (7.14)$$

4. The fitted value at  $x_0$  is given by  $\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0$ .

# Simulated Example

Local Regression



The blue curve is true  $f(x)$ , and the light orange curve is the local regression  $\hat{f}(x)$ . The orange points are local to the target point  $x_0$ , represented by the orange vertical line. The yellow bell-shape indicates weights assigned to each point. The fit  $\hat{f}(x_0)$  at  $x_0$  is obtained by fitting a weighted linear regression (orange line segment), and using the fitted value at  $x_0$  (orange solid dot) as the estimate  $\hat{f}(x_0)$ .

# Local Regression

- The size of the neighborhood (fraction  $s$  of training data) is a tuning parameter, which can be chosen by cross-validation.
- The weight of each point in the neighborhood needs to be specified.
- When we have two dimensional predictors  $X_1$  and  $X_2$ , we can simply use 2-dimensional neighborhoods, and fit bivariate linear regression models using the observations that are near each target point in 2-dimensional space.
- However, local regression can perform poorly if  $p$  is much larger than about 3 or 4 (the curse of dimensionality).
- $k$ -Nearest Neighbour is one of the most common local regression approaches. It corresponds to  $K(x_i, x_0) = \frac{1}{k}$  and  $\beta_1 = 0$ .

# Generalized Additive Models

- **Generalized additive models** (GAMs) provide a general framework for extending a standard linear model by allowing non-linear functions of each of the variables, while maintaining additivity,

$$y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}) + \epsilon_i.$$


- Each  $f_k$  can be linear, polynomials, step function, splines and local regression.
- Can be applied to classification problems.
  - ▶ Logistic regression:

$$\text{logit}(\mathbb{P}(Y_i = 1 \mid X_i = x_i)) = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}).$$

# Wage Data

Consider the wage data

$$\text{wage} = \beta_0 + f_1(\text{year}) + f_2(\text{age}) + f_3(\text{education}) + \epsilon.$$



The first two functions are natural splines in year and age. The third function is a step function, fit to the qualitative variable education.

# Pros and Cons of GAMs

- GAMs allow us to fit a non-linear  $f_j$  to each  $X_j$ , so that we can automatically model non-linear relationships that standard linear regression won't be able to capture.
- The non-linear fit can potentially improve prediction accuracy.
- Because the model is additive, we can still examine the effect of each  $X_j$  on  $Y$  individually while holding all of the other variables fixed.
- It avoids the curse of dimensionality by assuming additivity.
- However, GAMs fail to incorporate the interaction of variables.

# A Classical Local Approach: Nearest Neighbors

- Suppose we're given a new feature vector  $\mathbf{x} \in \mathbb{R}^p$  we consider classification.
- The idea: find the nearest feature vector to  $\mathbf{x}$  in the training set and use its label.
- Can formalize “nearest” in terms of the Euclidean distance

$$\|\mathbf{x}_i - \mathbf{x}_{i'}\|_2 = \sqrt{\sum_{j=1}^p (x_{ij} - x_{i'j})^2}$$

## Algorithm (1-NN):

1. Find example  $(\mathbf{x}_*, y_*)$  (from the stored training set) closest to  $\mathbf{x}$ .  
That is:


$$\mathbf{x}_* = \operatorname{argmin}_{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}} \text{distance}(\mathbf{x}_i, \mathbf{x})$$

2. Output  $y_*$  as the label




# Nearest Neighbors: Decision Boundaries

We can visualize the behavior in the classification setting using a **Voronoi diagram**.




# Nearest Neighbors: Decision Boundaries

**Decision boundary:** the boundary between regions of the feature space assigned to different categories.




# Nearest Neighbors: Decision Boundaries



Example: 2D decision boundary

# Nearest Neighbors



- Nearest neighbors **sensitive to noise or mis-labeled data** (“class noise”).
- Solution? Smooth by having  $k$  nearest neighbors vote


## Algorithm ( $k$ -NN):

1. Find  $k$  data points  $(\mathbf{x}_{(1)}, y_{(1)}), \dots, (\mathbf{x}_{(k)}, y_{(k)})$  closest to the test instance  $\mathbf{x}$
2. Classification output is majority class

$$\arg \max_{y \in C} \sum_{i=1}^k 1 \{y = y_{(i)}\}.$$


# $k$ -NN

$k=1$



# $k$ -NN

$k=15$



## Tradeoffs in choosing $k$

- Small  $k$ 
  - ▶ Good at capturing fine-grained patterns, more flexible decision boundary but high variance
  - ▶ Overfitting: may be sensitive to random idiosyncrasies in the training data.
- Large  $k$ 
  - ▶ Less flexible decision boundary and smaller variance
  - ▶ Underfitting: may fail to capture important regularities.
- Balancing  $k$ 
  - ▶ Optimal choice of  $k$  depends on number of data points  $n$ .
  - ▶ Nice theoretical properties if

$$k \rightarrow \infty, \quad \text{and} \quad \frac{k}{n} \rightarrow 0 \quad (\text{ESL 2.4}).$$

- ▶ Rule of thumb: choose  $k$  from  $[1, \sqrt{n}]$  via cross-validation!



# Pitfalls: The Curse of Dimensionality

- $k$ -NN suffers the curse of dimensionality!
  - ▶ In high dimensions, “most” points are approximately the same distance because they are far away from each other.
- Saving grace: some datasets (e.g. images) may have low **intrinsic dimension**, i.e. lie on or near a low-dimensional manifold.






Image credit: [https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make\\_swiss\\_roll.html](https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_swiss_roll.html)

The neighborhood structure depends on the intrinsic dimension.

# Pitfalls: Normalization

- Nearest neighbors can be sensitive to the ranges of different features.
- Often, the units are arbitrary:




- Simple fix: **standardize** each dimension to be zero mean and unit variance.
- Caution: depending on the problem, the scale might be important!

# Pitfalls: Computational Cost

- Computational cost for **training**: 0
- Computational cost for classifying **test data**, per data point (un-modified algorithm)
  - ▶ Calculate  $p$ -dimensional Euclidean distances with  $n$  data points:  
 $\mathcal{O}(np)$
  - ▶ Sort the distances:  $\mathcal{O}(n \log n)$
- This must be done for *each* test data point, which is very expensive by the standards of a learning algorithm!
- Need to store the entire dataset in memory!
- Tons of work has gone into algorithms and data structures for efficient nearest neighbors with high dimensions and/or large datasets.

# Example: Digit Classification

- Decent performance when lots of data







- Yann LeCunn – MNIST Digit Recognition
  - Handwritten digits
  - 28x28 pixel images:  $d = 784$
  - 60,000 training samples
  - 10,000 test samples
- Nearest neighbour is competitive

	Test Error Rate (%)
Linear classifier (1-layer NN)	12.0
K-nearest-neighbors, Euclidean	5.0
K-nearest-neighbors, Euclidean, deskewed	2.4
K-NN, Tangent Distance, 16x16	1.1
K-NN, shape context matching	0.67
1000 RBF + linear classifier	3.6
SVM deg 4 polynomial	1.1
2-layer NN, 300 hidden units	4.7
2-layer NN, 300 HU, [deskewing]	1.6
LeNet-5, [distortions]	0.8
Boosted LeNet-4, [distortions]	0.7

# Example: Digit Classification

- Changing the distance measure can really improve  $k$ -NN.
- Example: shape contexts for object recognition. In order to achieve invariance to image transformations, they tried to warp one image to match the other image.
  - ▶ Distance measure: average distance between corresponding points on *warped* images
- Achieved 0.63% error on MNIST, compared with 3% for Euclidean KNN.
- Competitive with the state of the art at the time, but required careful engineering.



[Belongie, Malik, and Puzicha, 2002. Shape matching and object recognition using shape contexts.]